





KubeCon



CloudNativeCon

China 2021

Virtual





KubeCon



CloudNativeCon

China 2021

Virtual

Best Practice: DNS Failure Observability and Diagnosis in Kubernetes

Yuning Xie
Software Engineer
Alibaba Cloud

DNS in Kubernetes

Why DNS is critical?

Built-in observabilities in CoreDNS

BPF-based Tools for diagnosis

An example of diagnostic procedure

DNS in Kubernetes



KubeCon




CloudNativeCon

Virtual

China 2021

We need DNS when:

- One microservice app talks to another
- Peer-to-peer comms in database cluster
- Monitoring, logging agents talks to API Server



Under the hood



KubeCon



CloudNativeCon

Virtual

China 2021

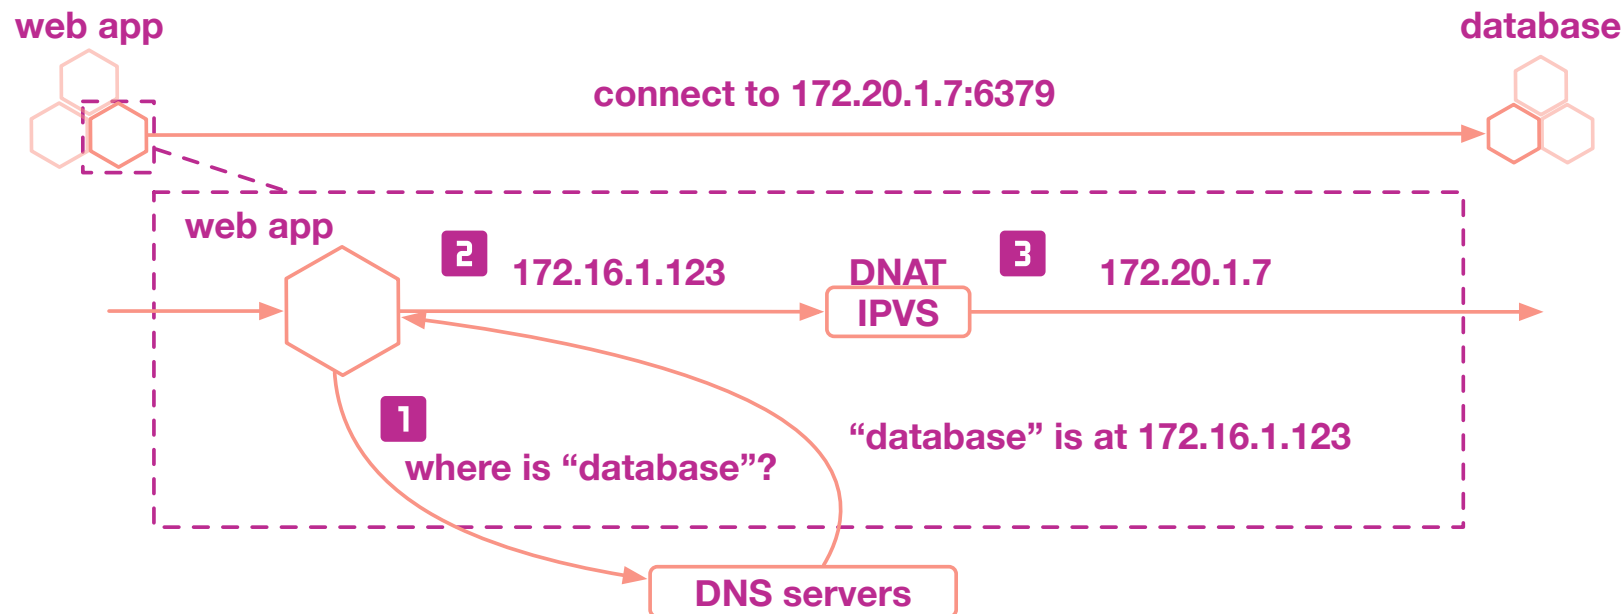
What it appears

- A TCP connection established.

```
[root@web-app-0 /]# telnet database 6379
Trying 172.21.7.76...
Connected to database.
Escape character is '^]'.
```

Under the hood

- App only knows "database" :6379
- App must know IP address first



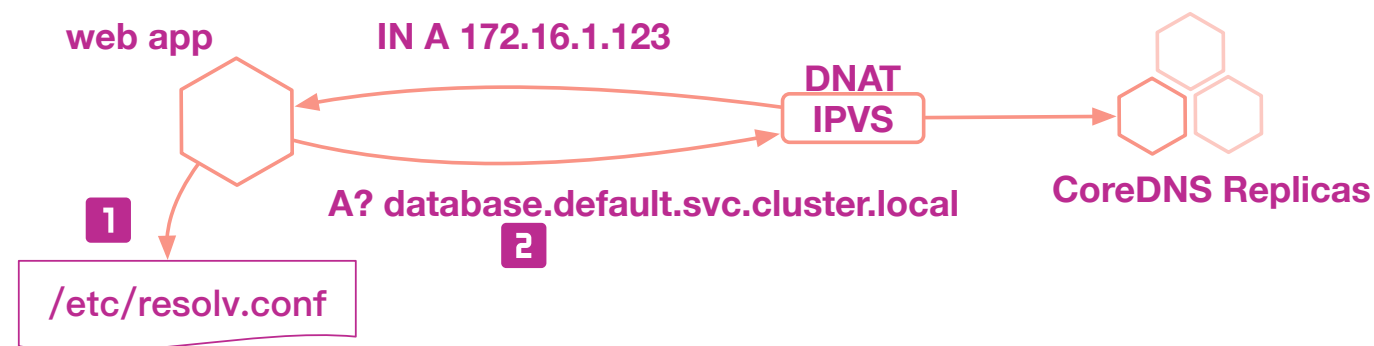
How it works

- Pods are configured to use in-cluster DNS servers
- The in-cluster DNS itself is a **Kubernetes Service**
- redundant search paths and ndots bring Service Discovery

Lots of DNS queries

- half of them are wasted AAAA queries
- more queries if domain is not local, e.g. querying A of www.aliyun.com will cost **8 queries**.

```
[root@web-app-0 /]# cat /etc/resolv.conf
nameserver 172.21.0.10
search default.svc.cluster.local svc.cluster.local cluster.local
options ndots:5
```



```
[root@web-app-0 ~]# tcpdump -nnn port 53
14:47:16.363701 IP 172.20.2.7.54526 > 172.21.0.10.53: 28204+ A? database.default.svc.cluster.local. (52)
14:47:16.363787 IP 172.20.2.7.54526 > 172.21.0.10.53: 2863+ AAAA? database.default.svc.cluster.local. (52)
14:47:16.364175 IP 172.21.0.10.53 > 172.20.2.7.54526: 28204*- 1/0/0 A 172.21.7.76 (102)
14:47:16.364184 IP 172.21.0.10.53 > 172.20.2.7.54526: 2863*- 0/1/0 (145)
```

DNS becomes fragile



CloudNativeCon

Virtual

China 2021

DNS itself is simple and reliable, but

- Almost every component needs DNS – **Huge Blast Radius**
- Each DNS query undergo too many components – **Vulnerable**
- For a single result of DNS query, app may trigger up to 8 queries – **High QPS**

DNS failures are becoming the top common issues in daily operations, but these failures' root causes vary.

DNS is probably the busiest protocol inside K8s.

CPU Starvation

Phenomenon: DNS delays under heavy load

Reason:

- DNS QPS is proportional to CPU limit

Mitigation: increase CPU limit or scale out CoreDNS replicas

Conntrack Table Entries Starvation

Phenomenon: DNS failure under heavy load

Reason:

- Each Service access will consume a conntrack entry
- If short-lived connection is widely used, Conntrack Table will soon be full
- DNS cannot be queried if Conntrack Table is full

Mitigation: switch to long-lived connections or increase Conntrack Table size

UDP is unreliable, from the day it was built.

Conntrack drops UDP packets when source port conflicts

Phenomenon: intermittent DNS delays

Reason:

- DNS itself is a Service in K8s
- Parallel requests to a Service reusing a same source port will leads to race condition
- Many resolvers lookup A and AAAA in parallel, e.g., musl in alpine

Mitigation: upgrade kernel to 4.19+, kubernetes issue#56903

IPVS drops UDP packets when source port conflicts

Phenomenon: intermittent DNS delays during CoreDNS rolling update or scaling in

Reason:

- IPVS does not expire connections after real server deleted
- When source port of an existing connection reused, UDP packets will be dropped

Mitigation: upgrade kernel to 5.9+, kubernetes issue#71514

Some legacy versions of CoreDNS have bugs.

CoreDNS plugins stop working

Phenomenon: intermittent lookup failure for Headless Service and external domains

Reason:

- Mishandle of Kubernetes Update events causing watch panics, and stop updating caches, coredns issue#3860
- Autopath stop working and return NXDOMAINs for external domains, coredns issue#4366

Mitigation: upgrade CoreDNS to latest



KubeCon



CloudNativeCon

Virtual

China 2021

DNS in Kubernetes


Built-in observabilities in CoreDNS

How do we observe these failure?

BPF-based Tools for diagnosis

An example of diagnostic procedure

Built-In Observabilities



Virtual

China 2021

Luckily, CoreDNS is a DNS server that chains a lot plugins, including many observability tools.

Observability	Plugin Name
Logging	log
	dump
	debug
	dnstap
Tracing	trace
Metrics	prometheus

Usage: log enables query logging to standard output.

Example:

```
[root@linux ~]# kubectl -n kube-system logs coredns-8fdd9787b-2bjvz --timestamps | grep database
2021-10-26T07:43:24.507437718Z [INFO] 172.20.2.7:48376 - 51559 "AAAA IN
database.default.svc.cluster.local. udp 52 false 512" NOERROR qr,aa,rd 145 0.000135929s
2021-10-26T07:43:24.507477682Z [INFO] 172.20.2.7:48376 - 30565 "A IN database.default.svc.cluster.local.
udp 52 false 512" NOERROR qr,aa,rd 102 0.000174301s
```


Note:

- Logs are printed on responding instead on receiving
- Timestamps can be fetched from "kubectl --timestamps"

Use Cases:

- Inspect a query's result at server-side
- Auditing cluster's outgoing domain queries, alert if any risky domains accessed

Loggings - log

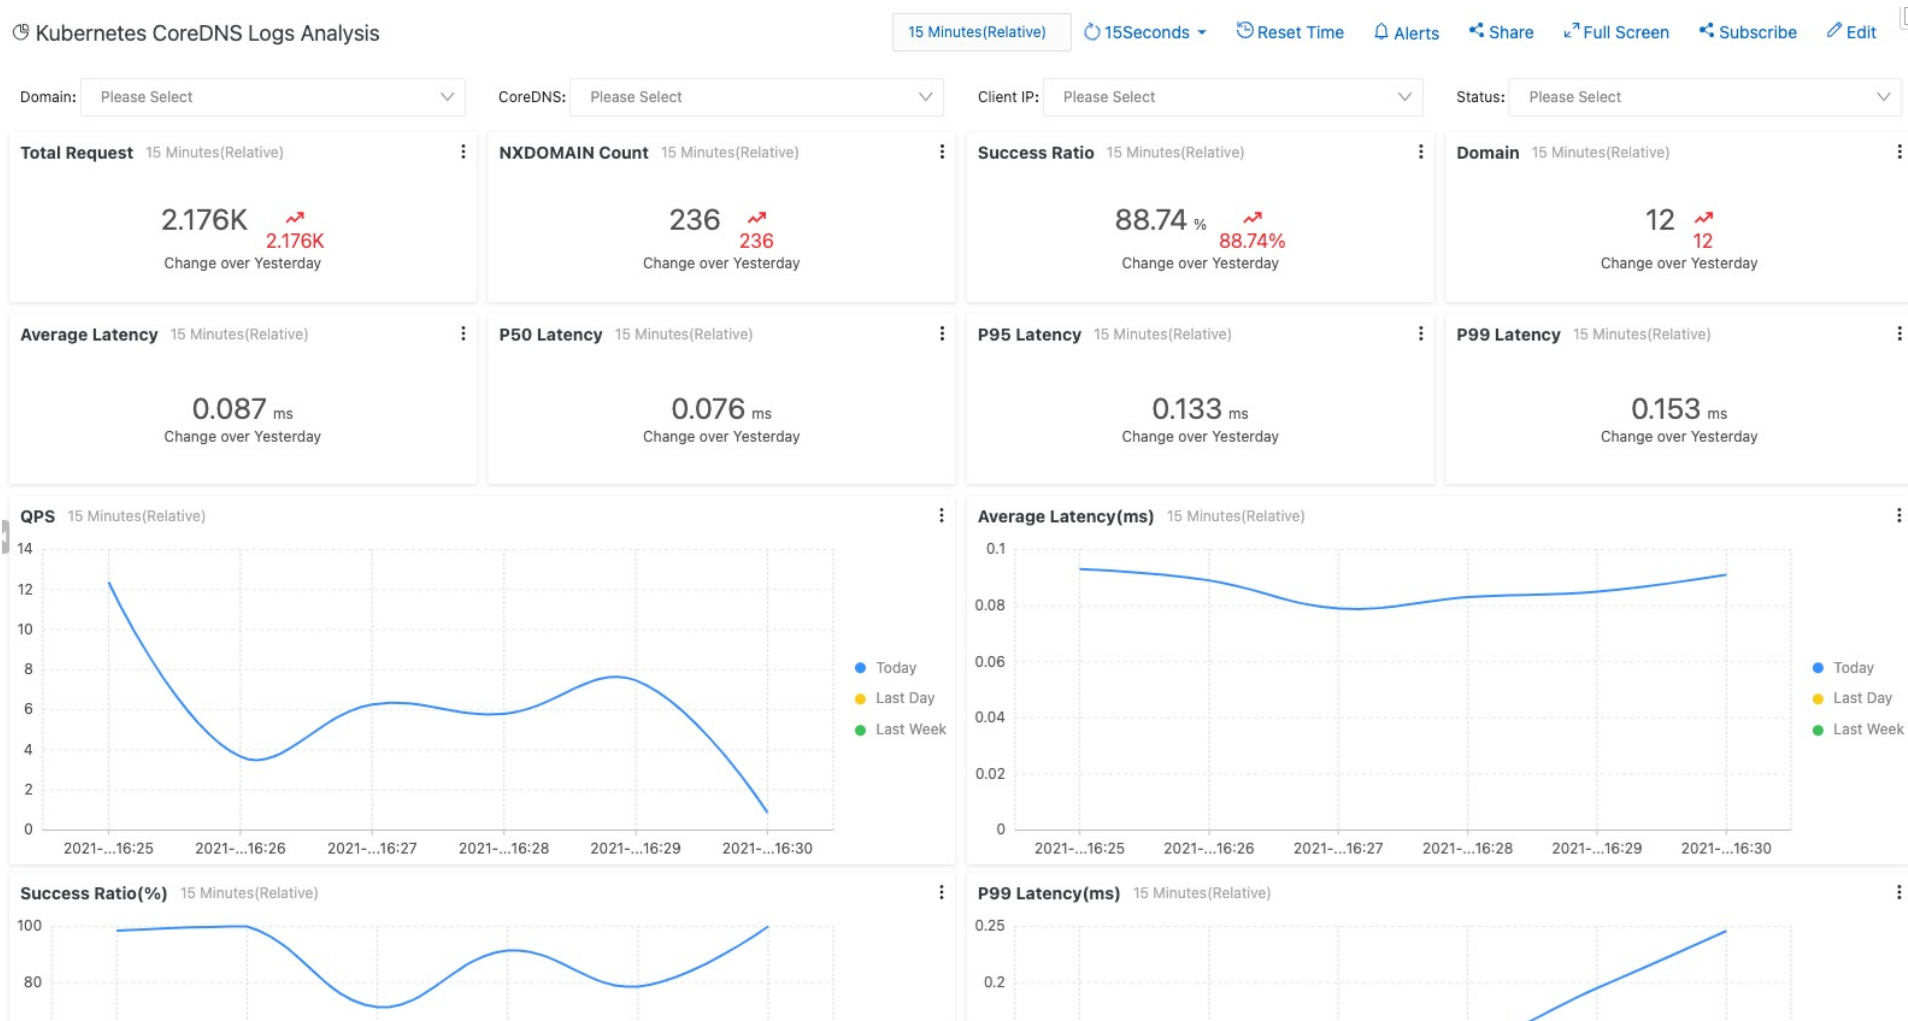


CloudNativeCon

Virtual

China 2021

Example: Auditing and analysis using Logging Dashboard



Loggings - dump



CloudNativeCon

Virtual

China 2021

Usage: dump dumps all incoming queries on standard output.

Example:

```
[root@linux ~]# kubectl -n kube-system logs coredns-8fdd9787b-2bjvz --timestamps | grep database
192.168.6.110 - 9858 A IN database.default.svc.cluster.local. udp 41973
```

Note:

- Dump is an external plugin of CoreDNS, must recompile CoreDNS to use
- Dump print queries on receiving

Use Cases:

- If you believe DNS queries arrived at CoreDNS but not logs shown, you should try this

Usage: debug disables the automatic recovery upon a crash so that you'll get a nice stack trace.

Example:

```
debug: 000000 00 0a 01 00 00 01 00 00 00 00 00 01 07 65 78 61
debug: 000010 6d 70 6c 65 05 6c 6f 63 61 6c 00 00 01 00 01 00
debug: 000020 00 29 10 00 00 00 80 00 00 00 00
debug: 00002a
```

Note:

- Probably not suitable for production environment

Use Cases:

- When upstream is sending strange replies, debug can output raw hex dump of DNS response

Loggings - dnstap



KubeCon



CloudNativeCon

Virtual

China 2021

Usage: dnstap enables logging to dnstap.

What is dnstap: dnstap is a flexible, **structured binary log format** for DNS software. It uses Protocol Buffers to encode events that occur inside DNS software in an implementation-neutral format.

How:


1. CoreDNS **send every DNS packets** to dnstap server
2. Dnstap server **store and analyze the packets**
3. Dnstap can alert on any abnormal queries

Advantages:


- **No performance degradation** on DNS servers
 - asynchronous I/O
 - buffered circular queue
 - no disk writes
- No text parsing
- **Extensible DNS message analysis**

Use Cases:

- Locating the issues, CoreDNS or upstream
- Alternative for tcpdump and stdout logs



Loggings - dnstap - how to analyze



CloudNativeCon

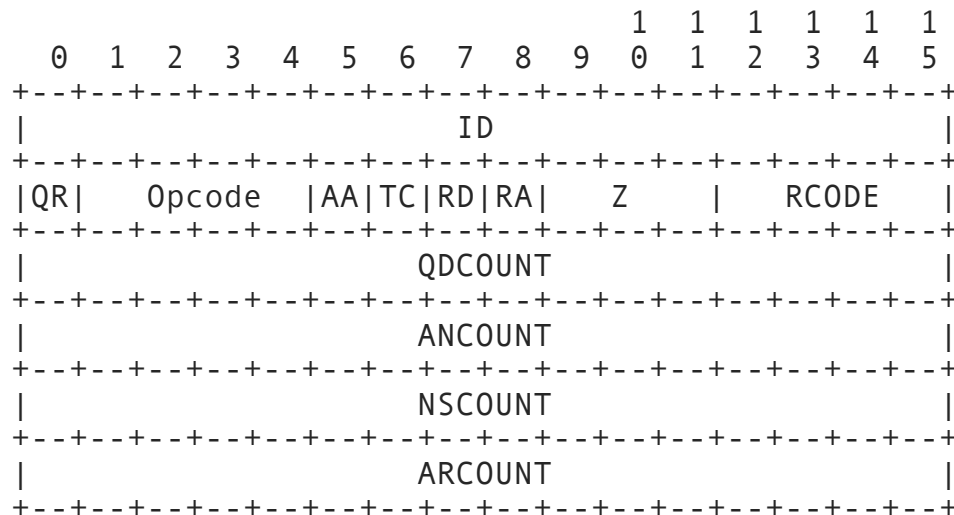
Virtual

China 2021

Response code(RCODE) is a signal

- when the receiving RCODE is bad, we should alert


DNS header format rfc1035



RCODE	Description
0	No error condition
1	Format error - The name server was unable to interpret the query.
2	Server failure - The name server was unable to process this query due to a problem with the name server.
3	Non-Existent Domain, or search path is incorrect
5	Query Refused

Loggings - dnstap - how to analyze

Dnstap - Message Types



Successful queries always come in pairs

- the dnstap receiver stores dnstap in LRU cache
- we can know which part is causing the errors by the contexts

Scenario	CLIENT_Q.	CLIENT_R.	FORWARDER_Q.	FORWARDER_R.
non-cached out of cluster domain query	yes	yes	yes	yes
cached out of cluster domain query	yes	yes	no	no
in-cluster domain query	yes	yes	no	no
cannot connect to upstream	yes	yes	yes	no

Tracing - trace



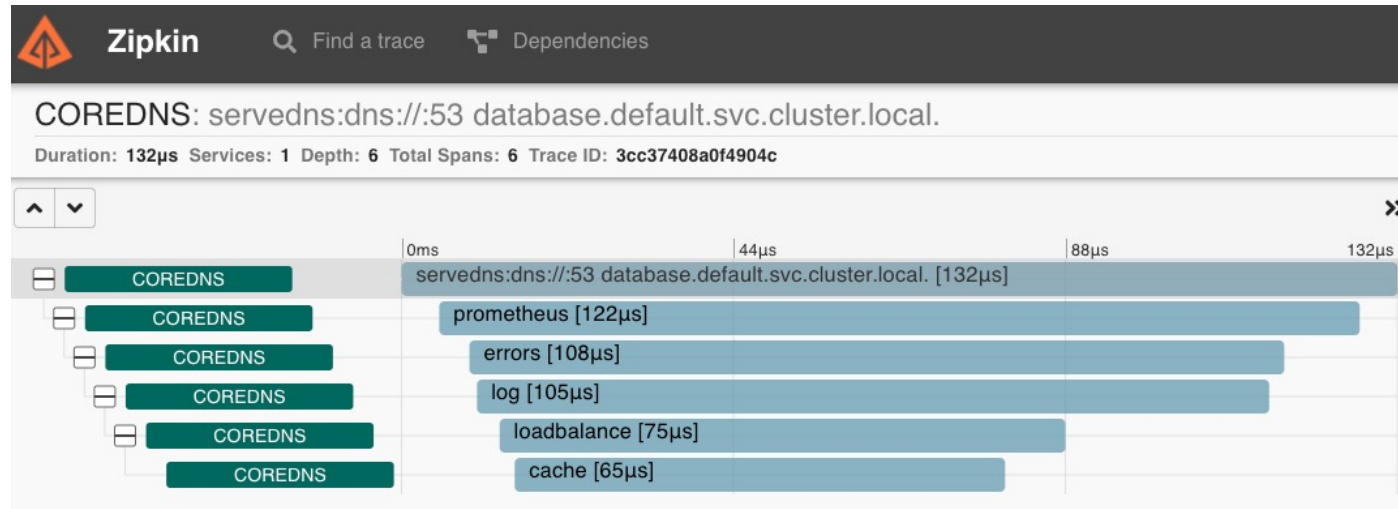
CloudNativeCon

Virtual

China 2021

Usage: trace enables OpenTracing-based tracing of DNS requests as they go through the plugin chain.

Example:



Note:

- Similar to dnstap, trace requires additional OpenTracing servers

Use Cases:

- To aid diagnosis on which plugin is slowing system down

Usage: prometheus enables Prometheus metrics.

Example:

```
[root@linux ~]# curl http://127.0.0.1:9153/metrics
# HELP coredns_cache_entries The number of elements in the cache.
# TYPE coredns_cache_entries gauge
coredns_cache_entries{server="dns://:53",type="denial"} 68
coredns_cache_entries{server="dns://:53",type="success"} 19
# HELP coredns_cache_hits_total The count of cache hits.
# TYPE coredns_cache_hits_total counter
coredns_cache_hits_total{server="dns://:53",type="denial"} 101302
coredns_cache_hits_total{server="dns://:53",type="success"} 96638
# HELP coredns_dns_request_duration_seconds Histogram of the time (in seconds) each request took.
# TYPE coredns_dns_request_duration_seconds histogram
coredns_dns_request_duration_seconds_bucket{server="dns://:53",type="A",zone=".",le="0.00025"} 99124
coredns_dns_request_duration_seconds_bucket{server="dns://:53",type="A",zone=".",le="0.0005"} 99312
coredns_dns_request_duration_seconds_bucket{server="dns://:53",type="A",zone=".",le="0.001"} 99382
```

Note:

- You will need Prometheus to scrape CoreDNS service at intervals

Use Cases:

- Provides statistical summaries about CoreDNS instances and DNS queries
- Alerts if any unexpected DNS RCODEs

Metrics - prometheus



KubeCon




CloudNativeCon

Virtual

China 2021

Example: Monitoring and alerting using Prometheus and Grafana





KubeCon



CloudNativeCon

Virtual

China 2021

DNS in Kubernetes

Built-in observabilities in CoreDNS

BPF-based Tools for diagnosis

What if failures happen at client-side?

An example of diagnostic procedure

Homemade BPF-based DNS Observer and Problem Diagnoser?

- UDP packets are not reliable by design
- BPF is great to locate the root cause

An example of locating DNS drops caused by IPVS drops UDP packets when source port conflicts and misconfiguration of iptables.

```
[root@linux ~]# bpftrace trace_dns_drop.bt
```

```
Attaching 8 probes...
```

```
Tracing DNS drops. Hit Ctrl-C to end.
```

TIME	PID	SKB	FUNC	LADDR:LPORT	RADDR:RPORT
15:06:24	203751	ffff8b8fe531a300	ip_vs_remote_request4	172.20.6.28:38888	172.21.0.10:53
15:06:28	203884	ffff8b8fe4b9ee00	ip_vs_local_request4	192.168.0.198:38888	172.21.0.10:53
15:07:04	204343	ffff8b8f85603d00	ipt_do_table	192.168.0.198:38888	172.20.6.51:53
15:07:05	204347	ffff8b8fa068fc00	ipt_do_table	172.20.6.28:38888	172.20.6.51:53
15:08:09	205378	ffff8b8f85603400	ip_vs_local_request4	192.168.0.198:38888	172.21.0.10:53
15:08:10	205491	ffff8b8f85708e00	ip_vs_remote_request4	172.20.6.28:38888	172.21.0.10:53

Code at <https://gist.github.com/xh4n3/61d8081b834d7e21bff723614e07777c>

Great tools:

- BCC gethostlatency.py <https://github.com/iovisor/bcc/blob/master/tools/gethostlatency.py>
- Packet, where are you? <https://github.com/cilium/pwru>



KubeCon



CloudNativeCon

Virtual

China 2021

DNS in Kubernetes

Built-in observabilities in CoreDNS

BPF-based Tools for diagnosis

An example of diagnostic procedure

How to deal with a real DNS incident?

An example of diagnostic procedure for DNS failures:



Confirm

- do not rush into conclusion too quickly
- config is desired and pods are running
- **no change recently**



Collect

- collect all possible info from coredns server, client, client's kernel
- **make dig tests from different places** to confirm the failure scope



Verify


- turn on CoreDNS's **log, metrics, trace plugins** to help diagnosis
- use **BPF tools or tcpdump** to verify where the packets dropped



Fix

- **perform the fix gradually**
- confirm result from metrics dashboard

Diagnostic Procedure



Virtual

China 2021

A couple of questions to help:

How often are failures?	Possible Cause	Next
Always	<ul style="list-style-type: none">• config error• network issue	<ul style="list-style-type: none">• check coredns log
Only on peak hours	<ul style="list-style-type: none">• resource limit on cpu, conntrack	<ul style="list-style-type: none">• check resource usage on nodes• check CoreDNS's metrics
Intermittently	<ul style="list-style-type: none">• conflicts issue caused by port reuse	<ul style="list-style-type: none">• use log plugin to locate cause first• check coredns pod status• check kernel version

How many nodes impacted?	Possible Cause	Next
Whole cluster	<ul style="list-style-type: none">• misconfig at server-side• kubernetes apiserver issue	<ul style="list-style-type: none">• check coredns config
Only a number of nodes	<ul style="list-style-type: none">• resource limit on cpu, conntrack	<ul style="list-style-type: none">• check network connectivity• check resource usage on nodes

Conclusion



KubeCon



CloudNativeCon

Virtual

China 2021

- DNS in Kubernetes is critical but fragile.
- CoreDNS has many observability plugins.
- BPF is helpful for locating in-kernel DNS failures.
- Follow best practice when set up CoreDNS in kubernetes, use NodeLocal DNSCache.

Thank You!